

INTRO TO DATA SCIENCE

LECTURE 10: SUPPORT VECTOR MACHINES

Paul Burkard

07/28/2015

LAST TIME:

- DECISION TREES**
- BUILDING DECISION TREES**
- OBJECTIVE SPLITTING FUNCTIONS**
- PREVENTING OVERFITTING IN DECISION TREES**

QUESTIONS?

I. SUPPORT VECTOR MACHINES

II. SOFT-MARGIN SVM

III. NONLINEAR SVM

HANDS-ON: SVM

I. SUPPORT VECTOR MACHINES

***Q:** What is a support vector machine?*

***A:** A **binary linear classifier** whose decision boundary is explicitly constructed to minimize generalization error.*

recall:

binary classifier — solves two-class problem

linear classifier — creates linear decision boundary (in 2d)

***Q:** How is the decision boundary derived?*

***A:** Using **geometric reasoning** (as opposed to the algebraic reasoning we've used to derive other classifiers).*

*The generalization error is equated with the geometric concept of **margin**, which is the region along the decision boundary that is free of data points.*




FIGURE 18-4. Two decision boundaries and their margins. Note that the vertical decision boundary has a wider margin than the other one. The arrows indicate the distance between the respective support vectors and the decision boundary.

***Q:** How is the decision boundary derived?*

***A:** Using **geometric reasoning** (as opposed to the algebraic reasoning we've used to derive other classifiers).*

*The goal of an SVM is to create the linear decision boundary with the **largest margin**. This is commonly called the **maximum margin hyperplane**.*

Q: *How is the decision boundary (**mmh**) derived?*

A: *By the **discriminant function**,*

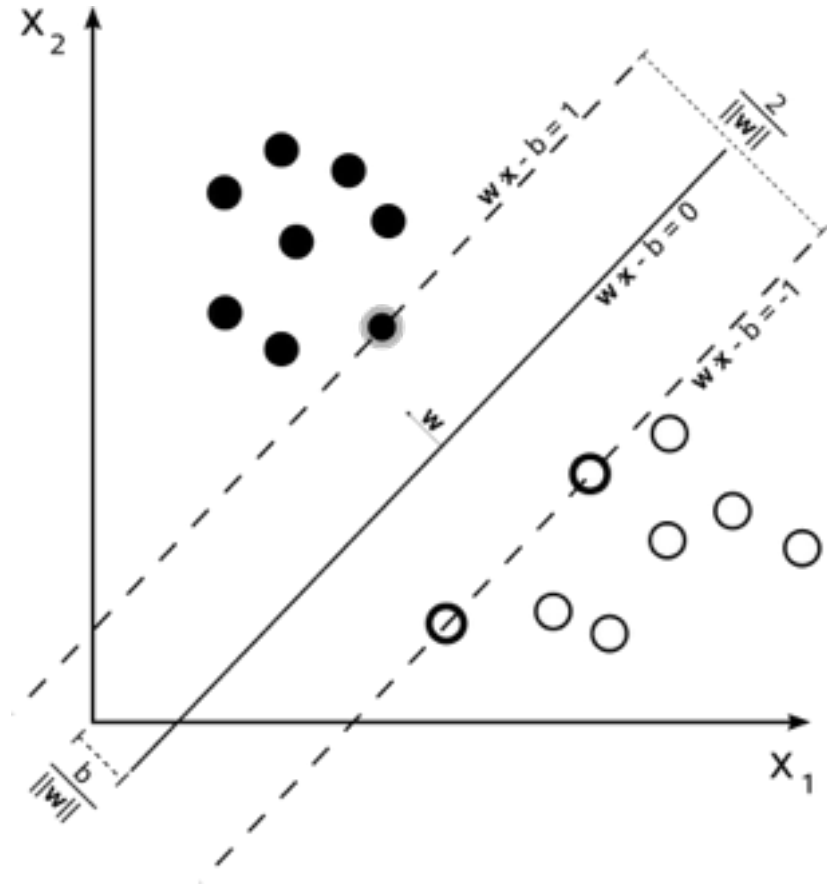
$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b.$$

such that w is the weight vector and b is the bias.

The sign of $f(x)$ determines the (binary) class label of a record x .


*How is the decision boundary (**mmh**) derived?*

- Any hyperplane can be written as the set of points where $\mathbf{w}^T \mathbf{x} - \mathbf{b} = 0$
- \mathbf{w} sets the plane's orientation (it's perpendicular to the plane)
- \mathbf{b} sets the offset from the origin
- Set the margin planes for each class such that $\mathbf{w}^T \mathbf{x} - \mathbf{b} = \pm 1$
- +1 for positive class, -1 for negative class




How is the decision boundary (mmh) derived?

- Thus, our goal of maximizing the margin width amounts to maximizing $2/\|w\|$
- To do this, we must minimize $\|w\|$
- Easier to minimize $\|w\|^2/2$
- We do this subject to the constraint:
 $y_i^*(w^T x_i - b) \geq 1$
- This is a **Quadratic Programming (QP)** optimization problem, yielding w and b



How is the decision boundary (mmh) derived?

- This is a **Quadratic Programming (QP)** optimization problem, yielding \mathbf{w} and \mathbf{b}
- Solving this QP problem yields a solution of the form:
$$\mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i$$
- Most of these alpha are 0
- Only the “**support vectors**” will be nonzero
- Thus only they contribute!



II. SOFT-MARGIN SVM

- *So far we've been dealing with what's called a **hard margin SVM** – what does this mean?*
- *We assume the data is perfectly separable by our hyperplane*
- *No classification instances will fall within the margin*
- *As we've seen, probably not realistic (especially in linear case)*
- *So how do we allow for some error/noise in our model?*

***Q:** How do we handle potentially inseparable data?*

***A:** By training a **soft margin SVM** rather than a hard margin*

soft margin – basically allows for a fuzzy boundary, where some proportion of elements may be misclassified in order to maintain a simpler boundary

Remember, simpler boundaries are probably more likely to generalize.

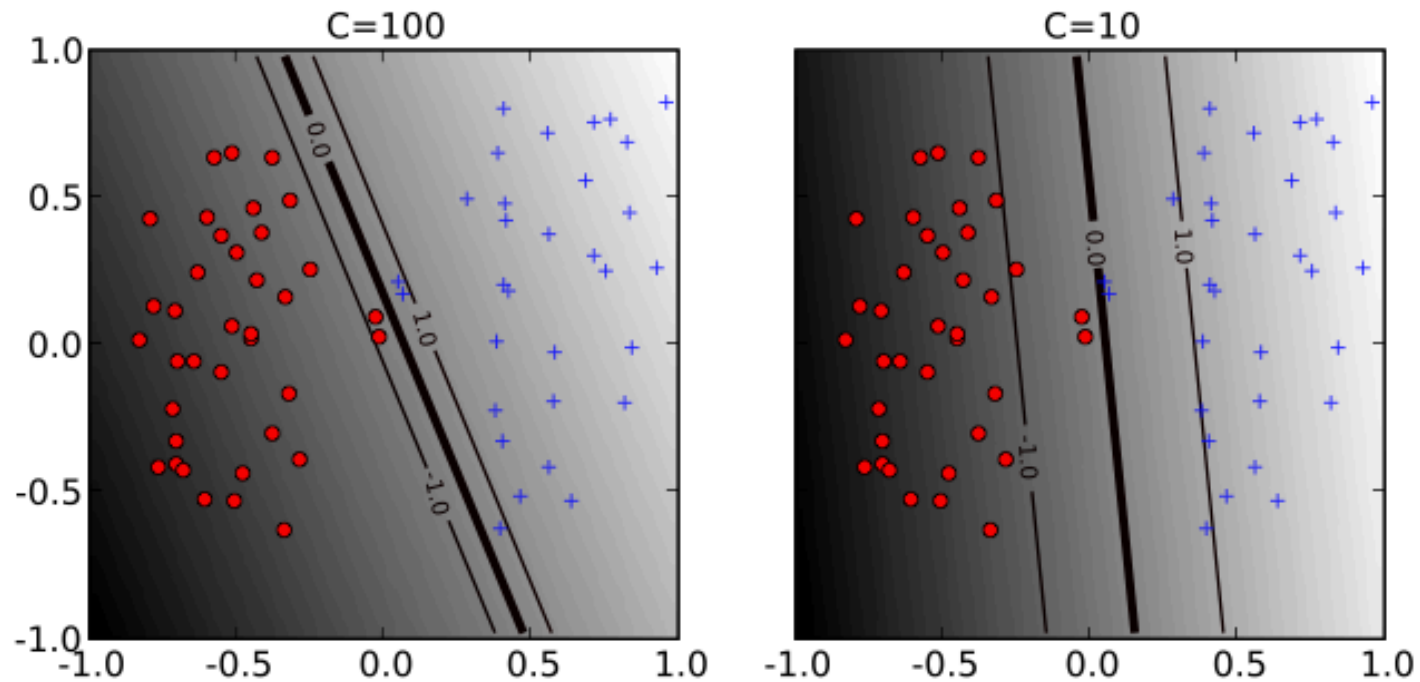
Q: How do we train a **soft margin SVM**?

A: By making use of **slack variables** that allow a proportion of instances to be misclassified

Slack variables ξ_i generalize the optimization problem to permit some misclassified training records (which come at a **cost C**).

The resulting soft margin QP problem is given by:

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \\ & \text{subject to:} && y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0. \end{aligned}$$



III. NONLINEAR SVM

Recall that our soft-margin SVM optimization function can be written as follows:

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \\ & \text{subject to:} && y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0. \end{aligned}$$

This can be rewritten in the following form:

$$\begin{aligned} & \underset{\alpha}{\text{maximize}} && \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i y_j \alpha_i \alpha_j \mathbf{x}_i^\top \mathbf{x}_j \\ & \text{subject to:} && \sum_{i=1}^n y_i \alpha_i = 0, \quad 0 \leq \alpha_i \leq C. \end{aligned}$$

NOTE

This is called the *dual formulation* of the optimization problem.

(reached via Lagrange multipliers)


*Notice that this expression depends only on **inner products** $\mathbf{x}_i^\top \mathbf{x}_j$*

The inner product (dot product) is an operation that takes two vectors and returns a real number.

The fact that we we can rewrite the optimization problem in terms of the inner product means that ***we don't actually have to do any calculations*** in the feature space K , ***just dot products***.

In particular, we can easily change K to be some other space K' .


Suppose we need a more complex classifier than a linear decision boundary allows.




Suppose we need a more complex classifier than a linear decision boundary allows.

One possibility is to add nonlinear combinations of features to the data, and then to create a linear decision boundary in the enhanced (higher-dimensional) feature space.

This ***linear*** decision boundary will be mapped to a ***nonlinear*** decision boundary in the original feature space.



original feature space K



higher-dim feature space K'

The logic of this approach is sound, but there are a few problems with this version.

In particular, this will not scale well, since it requires many high-dimensional calculations.

It will likely lead to more complexity (both modeling complexity and computational complexity) than we want.

Let's hang on to the logic of the previous example, namely:

- remap the feature vectors x_i into a higher-dimensional space K'
- create a linear decision boundary in K'
- back out the nonlinear decision boundary in K from the result

But we want to save ourselves the trouble of doing a lot of additional high-dimensional calculations. How can we do this?

Recall that our optimization problem depends on the features only through the inner product $\mathbf{x}_i^T \mathbf{x}_j$:

$$\begin{aligned} & \underset{\alpha}{\text{maximize}} && \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n y_i y_j \alpha_i \alpha_j \mathbf{x}_i^T \mathbf{x}_j \\ & \text{subject to:} && \sum_{i=1}^n y_i \alpha_i = 0, \quad 0 \leq \alpha_i \leq C. \end{aligned}$$

We can replace this inner product with a more general “***kernel function***” that has the same type of output as the inner product.

The upshot is that we can use a kernel function to *implicitly* train our model in a higher-dimensional feature space, *without* incurring additional computational complexity!

As long as the kernel function satisfies certain conditions, our conclusions above regarding the mmh continue to hold.

NOTE

These conditions are contained in a result called *Mercer's theorem*.

some popular kernels:

linear kernel

$$k(\mathbf{x}, \mathbf{x}') = \langle \mathbf{x}, \mathbf{x}' \rangle$$


polynomial kernel

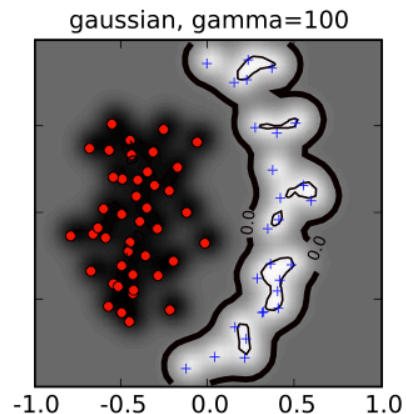
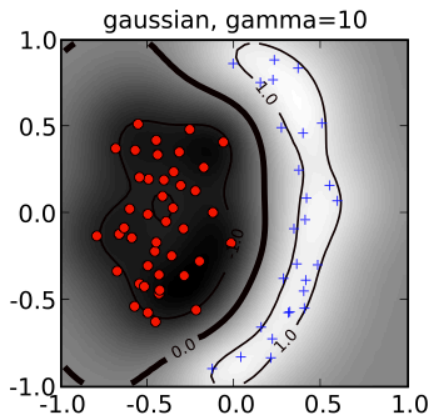

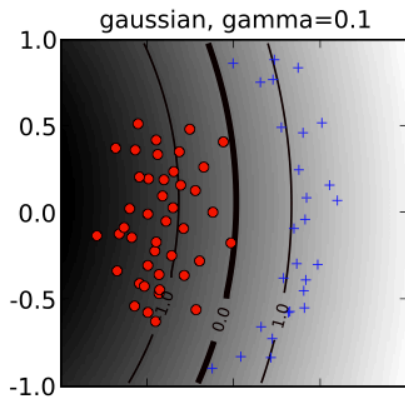
$$k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}' + 1)^d$$

Gaussian (RBF) kernel

$$k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2)$$

The **hyperparameters** d , γ affect the flexibility





SVMs (and kernel methods in general) are versatile, powerful, and popular techniques that can produce accurate results for a wide array of classification problems.

The main disadvantage of SVMs is the lack of intuition they produce. These models are truly black boxes!

INTRO TO DATA SCIENCE

HANDS-ON: SVM